

Importing necessary libraries

In [1]:

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
```

Importing datasets

In [2]:

```
path="tips.csv"
df=pd.read_csv(path)
```

Descriptive Statistics

In [3]:

```
# Preview data
df.head(10)
```

Out[3]:

	total_bill	tip	sex	smoker	day	time	size	price_per_person	Payer Name	CC
0	16.99	1.01	Female	No	Sun	Dinner	2	8.49	Christy Cunningham	356032510
1	10.34	1.66	Male	No	Sun	Dinner	3	3.45	Douglas Tucker	447807130
2	21.01	3.50	Male	No	Sun	Dinner	3	7.00	Travis Walters	601181210
3	23.68	3.31	Male	No	Sun	Dinner	2	11.84	Nathaniel Harris	467613760
4	24.59	3.61	Female	No	Sun	Dinner	4	6.15	Tonya Carter	483273260
5	25.29	4.71	Male	No	Sun	Dinner	4	6.32	Erik Smith	213140310
6	8.77	2.00	Male	No	Sun	Dinner	2	4.38	Kristopher Johnson	222372750
7	26.88	3.12	Male	No	Sun	Dinner	4	6.72	Robert Buck	351478500
8	15.04	1.96	Male	No	Sun	Dinner	2	7.52	Joseph Mcdonald	352286630
9	14.78	3.23	Male	No	Sun	Dinner	2	7.39	Jerome Abbott	353212450

In [4]:

```
# Dataset dimensions - (rows, columns)
df.shape
```

Out[4]:

(244, 11)

In [5]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 244 entries, 0 to 243
Data columns (total 11 columns):
#   Column                Non-Null Count  Dtype
---  -
0   total_bill            244 non-null   float64
1   tip                   244 non-null   float64
2   sex                   244 non-null   object
3   smoker                244 non-null   object
4   day                   244 non-null   object
5   time                  244 non-null   object
6   size                  244 non-null   int64
7   price_per_person      244 non-null   float64
8   Payer Name            244 non-null   object
9   CC Number             244 non-null   int64
10  Payment ID            244 non-null   object
dtypes: float64(3), int64(2), object(6)
memory usage: 21.1+ KB
```

In [6]:

```
df.isnull().sum()
#finding if our data has null values or not
```

Out[6]:

```
total_bill    0
tip           0
sex           0
smoker        0
day           0
time          0
size          0
price_per_person  0
Payer Name    0
CC Number     0
Payment ID    0
dtype: int64
```

In [16]:

```
df = df.drop(columns=['price_per_person', 'CC Number'])
```

In [17]:

```
df.describe()
```

Out[17]:

	total_bill	tip	size
count	244.000000	244.000000	244.000000
mean	19.785943	2.998279	2.569672
std	8.902412	1.383638	0.951100
min	3.070000	1.000000	1.000000
25%	13.347500	2.000000	2.000000
50%	17.795000	2.900000	2.000000
75%	24.127500	3.562500	3.000000
max	50.810000	10.000000	6.000000

In [18]:

```
correlation = df.corr()  
correlation
```

Out[18]:


	total_bill	tip	size
total_bill	1.000000	0.675734	0.598315
tip	0.675734	1.000000	0.489299
size	0.598315	0.489299	1.000000

Data Visualization

We will visualize correlation with heatmap, use count plots to see if the women or men come to the restaurant more than one another. Then let's see if the tip left really depends on the gender of the customer with box plot.


In [19]:

```
sns.heatmap(correlation)  
plt.show()
```




In [20]:

```
sns.catplot(x="sex",data=df,kind="count")  
plt.show()
```



In [21]:

```
sns.catplot(x="sex",y="tip",data=df,kind="box")  
plt.show()
```



Let's see how many days there are, then look at the relationship between the total bill paid and the tip, with respect to days.

In [22]:


```
df.day.unique()
```

Out[22]:

```
array(['Sun', 'Sat', 'Thur', 'Fri'], dtype=object)
```

In [23]:

```
sns.relplot(x="total_bill",y="tip",data=df,kind="scatter",hue="day")  
plt.show()
```



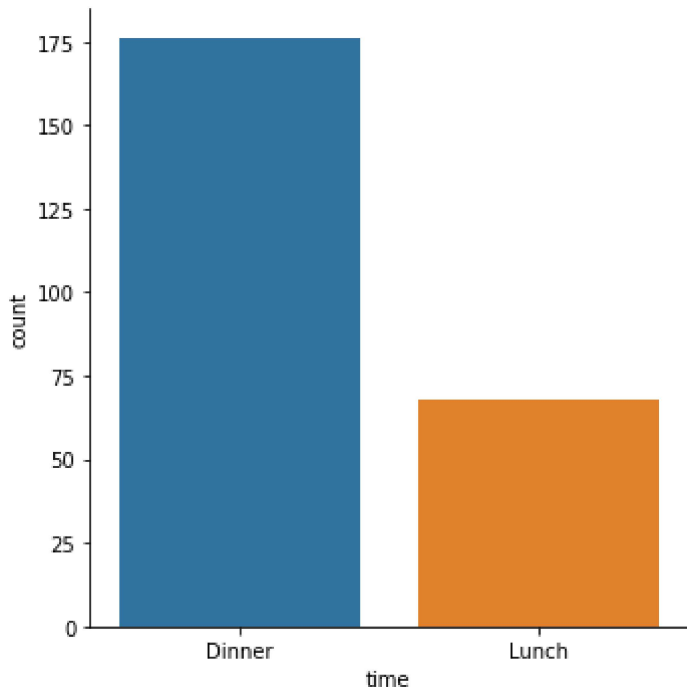
Let's see if people come over more in day time or in the evening.

In [24]:

```
sns.catplot(x="time", data=df, kind="count")  
plt.show
```

Out[24]:

```
<function matplotlib.pyplot.show(close=None, block=None)>
```



Let's see the correlation between the tip column with others.

In [25]:

```
df.corrwith(df["tip"])
```

Out[25]:

```
total_bill    0.675734  
tip           1.000000  
size          0.489299  
dtype: float64
```

Preprocessing

We want to predict the tip from other columns, therefore, We have to scale the numerical columns and encode categorical columns. For binary ones, you have to either use label encoding or one hot encode them, then drop duplicate ones.

In [26]:

```
columns_to_encode = ["sex", "smoker", "time"]
columns_to_scale = ["total_bill", "tip", "size"]
```

There are two unique values in sex, time and smoker columns and four values in day column, so we have to encode sex, time and smoker with label encoder and encode day with one hot encoder.

Now we are going to encode the day column with pandas' get_dummies method then directly append it to the main dataframe with pd.concat.

In [27]:

```
df = pd.concat([df, pd.get_dummies(df["day"], prefix="day")], axis=1)
df.head()
```

Out[27]:

	total_bill	tip	sex	smoker	day	time	size	Payer Name	Payment ID	day_Fri	day_Sat
0	16.99	1.01	Female	No	Sun	Dinner	2	Christy Cunningham	Sun2959	0	0
1	10.34	1.66	Male	No	Sun	Dinner	3	Douglas Tucker	Sun4608	0	0
2	21.01	3.50	Male	No	Sun	Dinner	3	Travis Walters	Sun4458	0	0
3	23.68	3.31	Male	No	Sun	Dinner	2	Nathaniel Harris	Sun5260	0	0
4	24.59	3.61	Female	No	Sun	Dinner	4	Tonya Carter	Sun2251	0	0

We will import Label Encoder from sklearn, instantiate it, and fit and transform the columns that We want to label, then convert the output array into dataframe and insert it to the original one.

In [29]:

```
from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()
df["sex"] = pd.DataFrame(le.fit_transform(df["sex"]))
```

In [30]:

```
df["time"] = pd.DataFrame(le.fit_transform(df["time"]))
df["smoker"] = pd.DataFrame(le.fit_transform(df["smoker"]))
```

Now we will import MinMaxScaler to normalize the numerical columns, and put them in another dataframe called scaled_columns, then drop the original columns and append them to the original dataframe.

In [31]:

```
from sklearn.preprocessing import MinMaxScaler
scaler = MinMaxScaler()
scaled_columns = pd.DataFrame(scaler.fit_transform(df[columns_to_scale]), columns=columns_to_scaled_columns)
scaled_columns.describe()
```

Out[31]:

	total_bill	tip	size
count	244.000000	244.000000	244.000000
mean	0.350145	0.222031	0.313934
std	0.186477	0.153738	0.190220
min	0.000000	0.000000	0.000000
25%	0.215281	0.111111	0.200000
50%	0.308442	0.211111	0.200000
75%	0.441087	0.284722	0.400000
max	1.000000	1.000000	1.000000

In [32]:

```
df.drop(["total_bill", "tip", "size", "day"], axis=1, inplace=True)
df = pd.concat([df, scaled_columns], axis=1)
```

Let's see the last version of the dataframe. Everything should be numeric.

In [33]:

```
df.head()
```

Out[33]:

	sex	smoker	time	Payer Name	Payment ID	day_Fri	day_Sat	day_Sun	day_Thur	total_bill
0	0	0	0	Christy Cunningham	Sun2959	0	0	1	0	0.291579
1	1	0	0	Douglas Tucker	Sun4608	0	0	1	0	0.152283
2	1	0	0	Travis Walters	Sun4458	0	0	1	0	0.375786
3	1	0	0	Nathaniel Harris	Sun5260	0	0	1	0	0.431713
4	0	0	0	Tonya Carter	Sun2251	0	0	1	0	0.450775

In [34]:

```
correlation = df.corr().abs()
plt.figure(figsize=(8,8))
sns.heatmap(correlation, annot=True)
plt.show()
```

